

# Cloud Computing Jumpstart

Gautham Pai

jnaapti

# Setting the Context

# How much data are we talking about really?

- ▶ 200 million Tweets per day – as of Jun 2011 (10s of GB of Tweets per day)
- ▶ Wikipedia dump – 31GB uncompressed
- ▶ Common Crawl data – 10s of TBs
- ▶ Tumblr – adding 3TB of new data everyday

# How much data are we talking about really?

- ▶ Google processes 25PB of data per day (as of 2011)
- ▶ Facebook – 4.6 million messages get sent every 20 minutes (as of Dec 2010)
- ▶ Youtube – 100+ hours of videos uploaded every minute (as of May 2013)

# We are dealing with a lot more data...

- ▶ Increase in the number of sensor devices
- ▶ Larger audience of users using our applications via the web and social networks results in increased data generation
- ▶ Cost of storage is falling – so we never discard any of the data

# Scrabulous – Case Study

- ▶ Scrabulous case study
  - ▶ Built by 2 young chaps from Kolkata
  - ▶ Both were in their early 20's when they built it
  - ▶ One was still in college.
  - ▶ 500,000 users daily – back in 2008, 25,000\$ in ad-revenues per month



Source: Wikipedia

Imagine this scenario

# We Have All It Takes


- ▶ We have access to a lot of the tools (software) that big corporations use for free
- ▶ We have access to a lot of the data for free
- ▶ We have computing power (hardware) available cheaply

This last point is because of the “Cloud Computing”  
revolution



# Questions to ask

- ▶ Ok, you have the resources
- ▶ You build a cool web application
- ▶ It is an overnight hit - can you handle it?
- ▶ What happens if the server has a disk crash?
- ▶ Can we prevent website outages in the account of hardware failures?



**Slashdot Effect**

# Looking for answers

What do technology companies like Google/Facebook/Twitter use to manage data? What challenges do they face in managing such huge volumes of data? How do they analyze such data?



Image Source: <http://opencompute.org/>


Making Systems Speak

# Processing Power over the Years

- ▶ Smartphones of today are more capable than desktops that you could buy a decade ago
- ▶ In 2002, you could buy a desktop with a single core of 400MHz.
- ▶ Today, we can get a dual-core 1GHz mobile phone for the same price

# Moore's Law


Microprocessor Transistor Counts 1971-2011 & Moore's Law



Source: Wikipedia

# Multicore

- ▶ We are reaching the limits of silicon
- ▶ Manufacturers are now packing more cores on a single chip
- ▶ Does a “kilo-core” processor sound weird?



Source: <http://www.flickr.com/photos/negativz/41549347/>

# Issue with Multicore

- ▶ Let's say this program takes 'x' seconds to run on a 1GHz processor

```
sum = 0
for i in range(10000000):
    sum += i
print sum
```

- ▶ How long does it take to run on:

- ▶ A 2GHz processor?
- ▶ A dual-core processor with 1GHz per core?

# Solution?

- ▶ Divide the work into smaller chunks that can be distributed across these cores

How about dividing the work across thousands of systems? Wouldn't that be even faster?!



# Distributed System


- ▶ A collection of computers which work together to achieve some task.
- ▶ The computers interact over a network.
- ▶ Google/Facebook/Twitter etc have tens of thousands of such systems working together on tasks such as crawling the web, coming up with friend recommendations, indexing your tweets etc.

# Working of a Website


- ▶ You enter a URL in the address bar of your browser
- ▶ You press ENTER
- ▶ You see a webpage

What happens behind the scenes?

# Client/Server Model



# TCP/IP




Source: Wikipedia

# TCP/IP – Low Level Details

- ▶ Tooling
  - ▶ DNS – dig
  - ▶ Finding your system's IP – ifconfig
  - ▶ Routing – traceroute
  - ▶ Finding socket information – netstat
  - ▶ You can use a tool called Wireshark to get packet information


# Client/Server Model – Separate DB Layer



# Problem 1 – Too Many Requests

- ▶ What if a thousand users access my server at the same time?
- ▶ If the server can handle 200 such requests in parallel in one second, what if I have 400 requests per second?
  - ▶ 1<sup>st</sup> second → 200 requests
  - ▶ 2<sup>nd</sup> second → 600 requests (200 are from the previous second)
- ▶ Results in server thrashing
- ▶ Solution: Load Balanced Setup

# Load Balancing






# Load Balancing

- ▶ Load balancing is a way of parallelizing processing across multiple machines
- ▶ The load balancer acts as a proxy that streams requests and responses between the client and the processing server.
- ▶ Eg: HAProxy
- ▶ Stateful and Stateless Architectures


# Problem 2 – Even More Requests

- ▶ What if the Load Balancer itself becomes the bottleneck?
- ▶ Solution:
  - ▶ Round Robin DNS
  - ▶ Building multiple independent clusters

# Clustered Setup



# Problem 3



# Virtualization

# Virtualization

- ▶ A way to use something “virtually”
- ▶ Eg: Running one operating system within another
  - ▶ The OS that runs the other OS is called the Host
  - ▶ The OS that runs within the other OS is the Guest

What if we have a way to manage such virtual instances  
via commands and API's?

# Virtualization Types

- ▶ Hardware Virtualization
- ▶ Application Virtualization
- ▶ Operating System Level Virtualization
- ▶ Others
  - ▶ Memory
  - ▶ Network
  - ▶ Storage

# Virtualization Techniques

- ▶ Virtual Machines (eg: KVM)
- ▶ Paravirtualization (eg: Xen)
- ▶ Containers (eg: LXC)



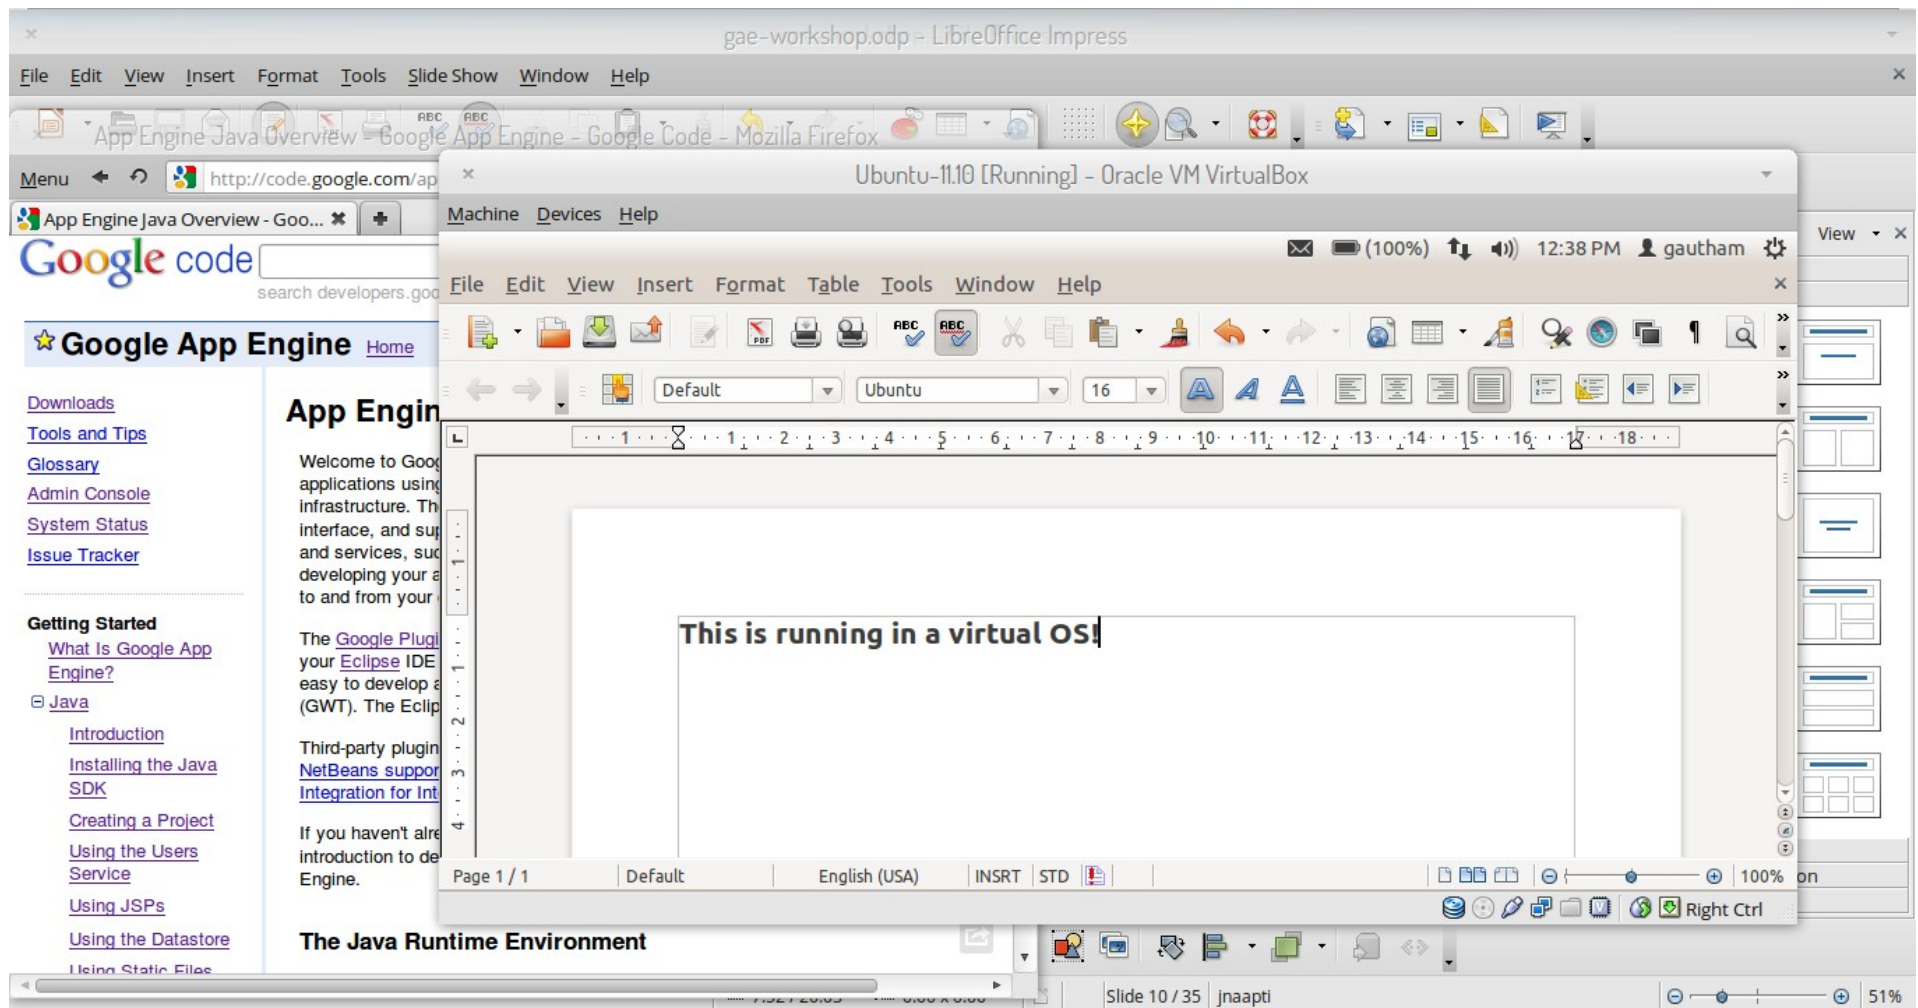
# Virtual Machine

- ▶ Process Virtual Machine (eg: application virtualization via JVM)
- ▶ System Virtual Machine (eg: KVM)

# Hypervisor

- ▶ Also called Virtual Machine Monitor
- ▶ A piece of computer software, firmware or hardware that creates and runs virtual machines
- ▶ Type-1 hypervisor or bare-metal hypervisor (eg: KVM)
- ▶ Type-2 hypervisor (eg: Virtualbox)

# Virtualization with VirtualBox – An Example



# Cloud Computing

# Cloud Computing

Cloud computing is the use of computing resources (hardware and software) that are delivered as a service over a network (typically the Internet).

- Wikipedia

# Have you used the cloud?

- ▶ If we talk about use of Cloud in terms of IaaS or PaaS, we may not have used it
- ▶ But if we talk about use of Cloud in terms of SaaS, there is a good possibility we have used it

# A Little Bit of History

- ▶ The pains of procuring systems back in early 2000...
- ▶ Virtualization comes into picture
- ▶ It's easy to provide resources via the Internet
- ▶ Cloud computing becomes a reality

# Then v/s Now

- ▶ Outlook v/s Gmail
- ▶ Local CRM v/s Salesforce
- ▶ MS Office v/s Google Docs
- ▶ Local storage v/s Dropbox
- ▶ Local enterprise storage v/s S3
- ▶ Expensive Server provisioning v/s EC2 and other AWS Services



# Benefits

- ▶ No upfront costs involved
- ▶ Lesser ongoing costs (lesser sysadmins required, no hardware failures headaches)
- ▶ Short time use of resources costs much lesser
- ▶ Meeting peak demands and bursts with pay as you go model

# Characteristics of Cloud Computing

- ▶ API based access
- ▶ Cost
- ▶ Device independence
- ▶ Virtualization
- ▶ Multitenancy

# API Based Access

- ▶ Being able to manage resources via APIs
- ▶ Case study – Boto
  - ▶ A Python interface to access AWS Services
  - ▶ <https://github.com/boto/boto>

# Multitenancy

- ▶ A single instance of a software is used to serve multiple “tenants”
- ▶ Advantages
  - ▶ Cost savings
  - ▶ Data mining/analytics
  - ▶ Ease of administration (release/maintenance/bug fixes etc)
- ▶ Disadvantage
  - ▶ Complexity of customization
- ▶ Case study: Google AppEngine Namespaces

IaaS/PaaS/SaaS

# Cloud Computing Service Models

- ▶ The evolutionary path
  - ▶ Software as a Service
  - ▶ Platform as a Service
  - ▶ Infrastructure as a Service

# Examples of IaaS, PaaS, SaaS

- ▶ Google AppEngine – PaaS
- ▶ Amazon AWS – IaaS and PaaS
- ▶ Rackspace – IaaS
- ▶ Heroku – PaaS
- ▶ Zoho applications – SaaS
- ▶ Gmail, Google Docs, Google Maps – SaaS

# Cloud Deployment Models

- ▶ Public Clouds
- ▶ Private Clouds
- ▶ Hybrid Clouds
  - ▶ Cloud Bursting



# Hosting your own Cloud

- ▶ Openstack
- ▶ Cloudstack
- ▶ Eucalyptus

# Networking Tools

# Networking Tools

- ▶ ifconfig
- ▶ dig
- ▶ ping
- ▶ traceroute
- ▶ netstat
- ▶ tcpdump
- ▶ resolv.conf
- ▶ ssh
- ▶ scp/rsync

# Web Engineering in the Cloud

# Web Engineering – The Components

- ▶ DNS
- ▶ Load Balancer
- ▶ Web Servers
- ▶ Application Servers
- ▶ Scaling web and application servers
- ▶ Data Stores/Cloud Storage
- ▶ Data analysis and the Hadoop ecosystem

# Web Engineering – Storage

- ▶ Relational Databases
- ▶ NoSQL Databases
- ▶ Scaling datastores

# Architecting for the Cloud

- ▶ A case study
- ▶ Understanding real scenarios
- ▶ Web Applications
- ▶ Media Storage
- ▶ Advertisement
- ▶ Any machine can go down anytime

# Architecting for the Cloud

- ▶ Avoiding single points of failures
- ▶ Simulate scalability by trying with true peak loads
- ▶ Chaos monkey
- ▶ Monitor and make sure errors don't occur again
- ▶ Think about continual automation and automated administration



# Introduction to Map/Reduce

# Examples of Web Scale Data Analysis

- ▶ Distributed Grep - Look for a pattern in all the Tweets
- ▶ Inverted Index Building - This is what is used by search engines
- ▶ Sentiment Analysis
- ▶ Competition Analysis
- ▶ Log Analysis

# Understanding the problem of Analysis

- ▶ Unlike in the case of retrieving data, in the case of analysis, we need to read through everything, but reads are slow in the disk.
- ▶ Let's see a simple math:
  - ▶ 1 Hard Disk read speed is 100MB/s
  - ▶ 100 Hard Disks read in parallel gives 10GB/s!
- ▶ Can we exploit this parallelism?

# The Coin Counting Example

- ▶ You have a sack full of coins, and you are asked to separate them into 1, 2, 5 and 10 Rs coins and tell how many of each are present.
- ▶ Now, let's say you have few sacks full of coins and it will take you a lot of time to count it yourself – so you call a few other people to help you out.
- ▶ Now, let's say there is few rooms full of coins (like in some large temples in India) – how will you count them?

# Coin Counting Problem – in depth

- ▶ You can't add more people to the same room – the room is already full.
- ▶ You can get a few more rooms, ask people to take some coins to the other room and then do the counting there, and come back with the coins and the final count.
- ▶ This will mean a lot of “traffic” in the corridor.
- ▶ So what's a better solution?

# A Possible Solution to the Coin Counting Problem

- ▶ Unload the coins in different rooms rather than in the same room.
- ▶ Then get workers in different rooms. With an increase in coins, increase the number of rooms and workers.
- ▶ Let the workers in each room work independently.

This is how Map/Reduce frameworks work

# Traditional Parallel Processing

- ▶ Use of threads, sharing data, synchronization
- ▶ Results in Deadlocks, Livelocks, Starvation etc
- ▶ Handling failures is complex

Parallel Programming is hard this way.

# Requirements from a parallel processing framework

- ▶ Higher level programming constructs – don't need to deal with sockets, threading, locking, sharing data etc
- ▶ Manage failures - if a task fails or a system breaks down, we want the framework to transparently manage it
- ▶ Recoverability - If a system fails, another system must be able to pick up its workload
- ▶ Replication – if a system fails, we don't lose data – the framework should replicate data in multiple nodes
- ▶ Scalability – Adding more compute nodes should help us increase the compute capacity



# Pulling data Or Pushing Computations?

- ▶ Pulling data for computation results in a bottleneck
- ▶ Every “database store” also has a “processor”.
- ▶ Instead of pulling the data for computation, can we think about pushing the computation out to where the data resides?
- ▶ Computation is in "bytes", may be a few MB of object code, that is still trivial compared to the data it works on

# MapReduce

- ▶ Concept introduced by Google in 2004
- ▶ Framework is inspired by map and reduce functions found in functional programming languages
- ▶ Hadoop is an opensource implementation of MapReduce

# MapReduce Frameworks

- ▶ Data is spread throughout machines before starting the task
- ▶ Computation is done in the nodes where data is stored
- ▶ Data is replicated in multiple machines to increase reliability
- ▶ Tasks are executed on multiple nodes just in case one of them is running slow

# Using the Common Crawl Data – A Case Study

- ▶ The dump is a few 10s of TBs in size
- ▶ Where/How do you download it?
- ▶ Answer: You don't need to download it
- ▶ Instead you push your computation to where the data exists, perform your computation and then only fetch results you are interested in!

Google AppEngine

# Google AppEngine

Google app engine

buzyp1@gmail.com | [My Account](#) | [Help](#) | [Sign out](#)

Not yet advertising on Google? Want to drive more customers to your app? [Get \\$100 credit for Google AdWords.](#)

Dismiss

Application: gautham-app-01 [High Replication] ▼

[Community Support](#) « [My Applications](#)

## Main

[Dashboard](#)

[Instances](#)

[Logs](#)

[Versions](#)

[Backends](#)

[Cron Jobs](#)

[Task Queues](#)

[Quota Details](#)

## Data

[Datastore Indexes](#)

[Datastore Viewer](#)

[Datastore Statistics](#)

[Blob Viewer](#)

[Prospective Search](#)

[Text Search](#)

[Datastore Admin](#)

[Memcache Viewer](#)

## Administration


[Application Settings](#)

Version: All Versions ▼

## Charts ?

Requests/Second ▼

6 hrs 12 hrs 24 hrs 2 days 4 days 7 days 14 days 30 days



## Instances ?

Number of Instances - [Details](#)

Average QPS

Average Latency

Average Memory

1 total

0.133

451.0 ms

52.6 MBytes

Billing Status: Free - [Settings](#)

Quotas reset every 24 hours. Next reset: 20 hrs ?

## Resource

## Usage

Frontend Instance Hours	<div></div>	0%	0.00 of 28.00 Instance Hours
Backend Instance Hours	<div></div>	0%	0.00 of 9.00 Instance Hours
Datastore Stored Data	<div></div>	0%	0.00 of 1.00 GBytes

# Google App Engine

- ▶ Free to try out
- ▶ Supports Java, Python and Go languages
- ▶ Does not provide direct access to the system running the application
- ▶ Instead allows us to work with application via commands and an admin-console
- ▶ You can have your first deployed app in minutes!

# Google AppEngine – Handson

- ▶ Create an application in the dashboard
- ▶ Create the webapp2 Request Handler
- ▶ Test your application locally
- ▶ Upload your application



# Google App Engine contd...

- ▶ Provides several useful building blocks that we will need to build applications.

Example:

- ▶ Data Storage
- ▶ User Management
- ▶ Mail Services
- ▶ Caching
- ▶ XMPP
- ▶ Image Manipulation
- ▶ Task Queues
- ▶ ...

OpenStack

# OpenStack

- ▶ Compute - Nova
- ▶ Object Storage - Swift
- ▶ Block Storage - Cinder
- ▶ Image Management - Glance
- ▶ Networking - Quantum
- ▶ Dashboard - Horizon
- ▶ APIs – Python and Shell based access, AWS Compatible

# Hypervisors Supported

- ▶ KVM
- ▶ LXC
- ▶ QEMU
- ▶ UML
- ▶ VMware vSphere
- ▶ Xen
- ▶ PowerVM (IBM)
- ▶ Hyper-V (Microsoft)
- ▶ Bare Metal

# Devstack


- ▶ A documented shell script to build complete OpenStack development environments.
- ▶ Very simple to install and try out OpenStack
- ▶ You can install the cloud in the cloud with this to get started!


# Amazon Web Services


# AWS Console


AWS Management Console › Amazon EC2


Help ▾


 **AWS**  
Elastic Beanstalk


 **Amazon S3**


 **Amazon EC2**


 **Amazon VPC**


 **Amazon CloudWatch**


 **Amazon Elastic MapReduce**


 **Amazon CloudFront**


 **AWS CloudFormation**


 **Amazon RDS**


 **Amazon ElastiCache**


 **Amazon SQS**


 **AWS IAM**


 **Amazon SNS**

 **Amazon SES**

 **Amazon Route 53**

 **Amazon DynamoDB**

 **AWS Storage Gateway**

 **Amazon SWF**

**Navigation**

**Region:**  
US East (Virginia) ▾

**EC2 Dashboard**  
Scheduled Events

INSTANCES

Instances  
Spot Requests  
Reserved Instances

IMAGES

AMIs  
Bundle Tasks

ELASTIC BLOCK STORE

Volumes  
Snapshots

NETWORK & SECURITY

Security Groups  
Elastic IPs  
Placement Groups  
Load Balancers  
Key Pairs  
Network Interfaces

**Amazon EC2 Console Dashboard**

**Getting Started**

To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 instance.

[Launch Instance](#)

Note: Your instances will launch in the US East (Virginia) region.

**Service Health**

Service Status

Current Status	Details
✓ Amazon EC2 (US East - N. Virginia)	Service is operating normally <a href="#">View complete service health details</a>

Availability Zone Status

Current Status	Details
✓ us-east-1a	Availability zone is operating normally
✓ us-east-1b	Availability zone is operating

**My Resources**

You are using the following Amazon EC2 resources in the US East (Virginia) region: [Refresh](#)

1 Running Instance

1 Elastic IP

1 EBS Volume

1 EBS Snapshot

1 Key Pair

0 Load Balancers

0 Placement Groups

1 Security Group

**Scheduled Events**

✓ US East (Virginia): **No events** [Refresh](#)

**Related Links**

[Getting Started Guide](#)  
[Documentation](#)  
[All EC2 Resources](#)  
[Forums](#)  
[Feedback](#)

© 2008 - 2012, Amazon Web Services LLC or its affiliates. All rights reserved. | [Feedback](#) | [Support](#) | [Privacy Policy](#) | [Terms of Use](#) | An [amazon.com](#) company

# Amazon Web Services (some popular services)

- ▶ Amazon Elastic Compute Cloud (EC2)
- ▶ Amazon Simple Storage Service (S3)
- ▶ Elastic Block Storage (EBS)
- ▶ Elastic Load Balancing (ELB)
- ▶ Amazon Relational Database Service (RDS)
- ▶ Amazon DynamoDB
- ▶ Auto Scaling
- ▶ Amazon ElastiCache



# AWS Management Console

- ▶ Introduction to the AWS Management Console
- ▶ Working with various AWS Services
  - ▶ EC2
  - ▶ S3
  - ▶ RDS
  - ▶ ...

# Regions and Availability Zones

- ▶ Regions and Availability Zones
  - ▶ Each region is in a different geographic area
  - ▶ Each Availability Zone is completely isolated from the other and connected to others via low latency links

# EC2 Purchasing Models

- ▶ On-demand instances
- ▶ Reserved instances
- ▶ Spot instances

# EC2 Instances

- ▶ Selecting an AMI
- ▶ Choosing the Instance Type
- ▶ Choosing the Availability Zone
- ▶ Choose on-demand or spot
- ▶ Enable Cloudwatch monitoring (if required)

# EC2 Instance

- ▶ Choose the disk
- ▶ Tag your instance
- ▶ Associate a key-pair
- ▶ Configure the security group
- ▶ Review and Launch
- ▶ Associating an Elastic IP with the instance

# AWS – A Few Case Studies

- ▶ Building web applications in AWS
- ▶ Ad serving via AWS
- ▶ Storing and processing media (sound/video/images) in AWS

# Resources

- ▶ KVM
- ▶ Virtual Machine Manager
- ▶ QEMU Images
- ▶ AWS Documentation
- ▶ AWS Reference Architecture
- ▶ Google AppEngine Documentation